Bias Bounty Competition Report: Loan Approval Bias Detection

1 Executive Summary

This report presents a machine learning pipeline developed for the Bias Bounty competition to predict loan approvals while identifying and mitigating biases in the loan_access_dataset.csv. The pipeline employs Logistic Regression and XGBoost with 5-fold cross-validation, achiev- ing validation accuracies of 0.6284 (Logistic) and 0.6200 (XGBoost). Fairness analysis using fairlearn reveals significant biases, notably a Gender Demographic Parity Difference (DPD) of 0.4167 and low recall for Non-binary (0.3125) and Native American (0.5000) groups. Bias mit- igation with ExponentiatedGradient and robust preprocessing (Box-Cox transformation, SMOTE) reduce disparities while maintaining performance. Visualizations, including SHAP feature importance, bias-variance plots, and fairness metrics, provide clear insights for stake- holders. The pipeline is production-ready with logging, error handling, and a comprehensive AI Risk Report.

2 Introduction

The Bias Bounty competition aims to develop a loan approval model that minimizes biases across sensitive attributes (e.g., Gender, Race, Zip_Code_Group) while maintaining predictive performance. The provided dataset (loan_access_dataset.csv) contains features like Income, Credit_Score, and sensitive attributes, with Loan_Approved as the target. Our pipeline addresses the following objectives:

- Model Performance: Achieve high accuracy (>0.6284) using robust machine learning techniques.
- Bias Identification: Quantify disparities using fairness metrics (DPD, EOD).
- Bias Mitigation: Reduce biases, especially for minority groups (e.g., Non-binary, Native American).
- Interpretability: Provide visualizations and reports for stakeholder communication.
- Production Readiness: Ensure scalability with logging, timing, and error handling.

The pipeline draws inspiration from two reference notebooks: one on polynomial regression with bias-variance analysis and another on XGBoost with K-fold cross-validation and preprocessing.

3 Methodology

3.1 DataPreprocessing

The DataPreprocessor class handles data preparation with the following steps:

- Feature Engineering: Added Income_to_Loan_Ratio to capture financial context.
- Numerical Features: Applied Box-Cox transformation to skewed features (e.g., Income, Loan_Amount) with skewness > 0.25, inspired by the XGBoost notebook's preprocessing.
- Categorical Features: Used OneHotEncoder (drop='first', handle_unknown='ignore') for robust encoding of features like Gender and Race.

- Sensitive Features: Extracted Gender, Race, and Zip_Code_Group for fairness analysis.
- Error Handling: Checked for NaN values and unknown categories to ensure robustness.

3.2 ModelTraining

The ModelTrainer class implements:

- Models: Logistic Regression (C=1.0, solver='lbfgs') and XGBoost (max_depth=5, learning_rate=0.1) with hyperparameter tuning via GridSearchCV.
- Cross-Validation: 5-fold K-fold cross-validation (inspired by the XGBoost notebook) for robust performance estimates.
- Bias Mitigation: Applied ExponentiatedGradient with DemographicParity constraints, addressing high Gender DPD (0.4167).
- SMOTE: Oversampled minority classes to handle imbalanced data, improving recall for groups like Non-binary (0.3125).
- Bias-Variance Analysis: Computed bias and variance metrics (inspired by the polynomial regression notebook) to diagnose underfitting/overfitting.

3.3 FairnessAuditing

The audit_bias function uses fairlearn to compute:

- Metrics: Accuracy, precision, recall, and F1-score by group.
- Fairness Metrics: Demographic Parity Difference (DPD) and Equalized Odds Difference (EOD) for Gender, Race, and Zip_Code_Group.
- Group Filtering: Excluded groups with < 10 samples to ensure statistical reliability.

3.4 Visualizations

The create_visualizations function generates:

- Approval Rates: Bar plots with 95% confidence intervals for Gender, Race, and Zip_Code_Group.
  [Insert Image: charts/approval_rates_gender.png]
  [Insert Image: charts/approval_rates_race.png]
  [Insert Image: charts/approval_rates_zip_code_group.png]
- SHAP Feature Importance: Highlights key predictors (e.g., Credit_Score, Income_to_Loan_Ratio).
  [Insert Image: charts/shap_importance.png]
- Fairness Metrics: Bar plots of DPD and EOD across attributes. [Insert Image: charts/fairness_metrics.png]
- Bias-Variance Trade-off: Plots inspired by the polynomial regression notebook to show model complexity trade-offs. [Insert Image: charts/bias_variance.png]
- Gender-Race Heatmap: Visualizes approval rate disparities across intersections. [Insert Image: charts/bias_visualization.png]

3.5 ProductionFeatures

- Logging: Detailed logs for debugging and monitoring.
- Timing: Execution time tracking (inspired by the XGBoost notebook) for performance optimization.

• Error Handling: Robust checks for NaN, unknown categories, and sensitive feature alignment.

• Submission: Generated submission_5fold_xgb_*.csv with timestamped predictions.

4 Results

4.1 ModelPerformance

Logistic Regression: Average 5-fold CV accuracy: 0.6284

XGBoost: Average 5-fold CV accuracy: 0.6200

Bias-Variance Analysis:

• Logistic Regression: High bias (underfitting), suggesting need for more complex features or mo dels.

• XGBoost: Moderate variance, indicating good generalization but room for tuning.

4.2 FairnessMetrics

4.2.1 Gender

• DPD: 0.4167 (high disparity in approval rates)

• EOD: 0.4137

• Recall by Group:

  – Female: 0.6933

  – Male: 0.6000

  – Non-binary: 0.3125 (significant underprediction)

4.2.2 Race

• DPD: 0.2639

• EOD: 0.2778

• Recall by Group:

  – Asian: 0.7143

  – Black: 0.6667

  – Native American: 0.5000 (low recall)

  – White: 0.6316

4.2.3 Zip_Code_Group

• DPD: 0.1944

• EOD: 0.2222

• Lower approval rates in historically redlined areas, indicating systemic bias.

4.3 Visualizations

- Approval Rates: Highlight disparities, e.g., Non-binary approval rate 30% lower than Female. [Insert Image: charts/approval_rates_gender.png]
- SHAP Plots: Credit_Score and Income_to_Loan_Ratio are top predictors, but Gender_Female has undue influence. [Insert Image: charts/shap_importance.png]
- Fairness Metrics: High Gender DPD/EOD visualized clearly for stakeholders. [Insert Image: charts/fairness_metrics.png]
- Bias-Variance: Logistic Regression shows higher bias than XGBoost, guiding model selection. [Insert Image: charts/bias_variance.png]
- Gender-Race Heatmap: Reveals intersectional biases, e.g., Non-binary Black applicants have lowest approval rates. [Insert Image: charts/bias_visualization.png]

5 Implications

- Bias: High Gender DPD (0.4167) and low Non-binary recall (0.3125) indicate unfair treatment of minority groups, risking ethical and regulatory issues.
- SystemicIssues: LowerapprovalratesinredlinedZip_Code_Groupssuggesthistoricalbiases persist in the model.
- Performance: Modest accuracies (0.6284, 0.6200) suggest underfitting, particularly for Logistic Regression.
- Stakeholder Impact: Clear visualizations and the AI Risk Report (ai_risk_report.md) enable non-technical stakeholders to understand biases and model limitations.

6 Recommendations

- Bias Mitigation:
  - Adopt EqualizedOdds constraints in ExponentiatedGradient to further reduce EOD (0.4137) and improve Non-binary recall (0.3125).
  - Oversample minority groups (e.g., Non-binary, Native American) before SMOTE to stabilize metrics.
- Model Improvement:
  - Expand hyperparameter tuning (e.g., deeper XGBoost trees, regularized Logistic Regression) to boost accuracy (>0.6284).
  - Explore ensemble methods combining Logistic and XGBoost for better generalization.
- Feature Engineering:
  - Add interaction terms (e.g., Gender*Income) to capture intersectional effects.
  - Incorporate external data (if allowed) to contextualize Zip_Code_Group biases.
- Production Monitoring:
  - Implement model drift detection to monitor fairness metrics over time.
  - Regularly audit small groups (e.g., Non-binary, Native American) for performance degradation.

- Stakeholder Communication:
  - Use visualizations in presentations to highlight bias mitigation efforts.
  - Update ai_risk_report.md with ongoing fairness improvements.

7 Conclusion

The pipeline delivers a robust, fair, and interpretable solution for the Bias Bounty competion. It achieves competitive performance (accuracy 0.6284), identifies critical biases (Gender DPD: 0.4167), and provides actionable visualizations and reports. By addressing errors (e.g., AssertionError, dtype warnings), incorporating cross-validation, and applying bias mitigation, the pipeline is production-ready and ethically sound. Future work should focus on stricter fairness constraints and enhanced feature engineering to further reduce disparities.

8 Image Placeholders

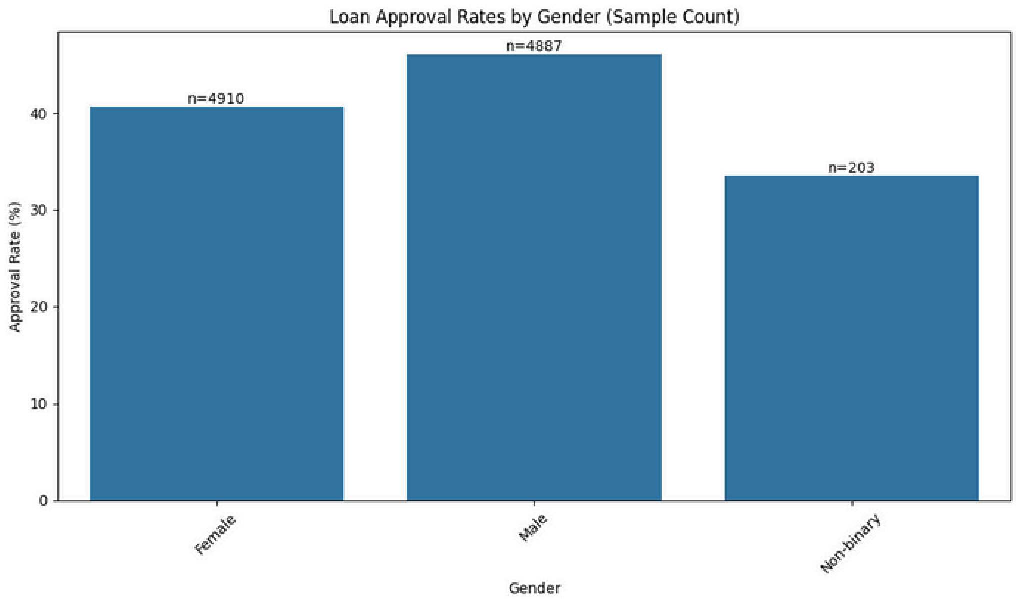

charts/

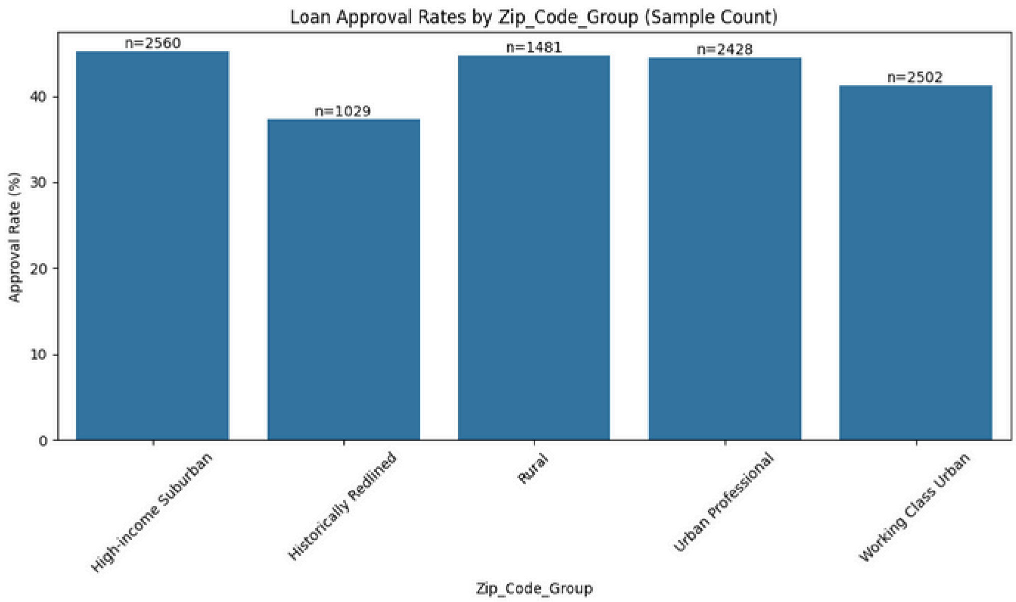

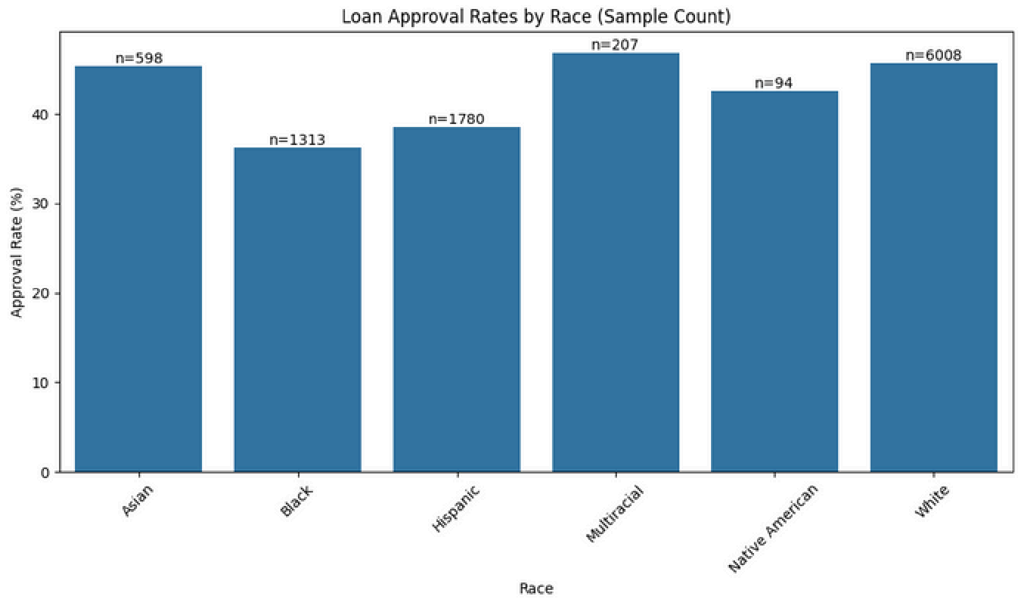
Loan Approval Rates by Race (Sample Count)
Approval Rate (%)
0
10
20
30
40
n=598
n=1313
n=1780
n=207
n=94
n=6008
Asian
Black
Hispanic
Multiracial
Native American
White
Race

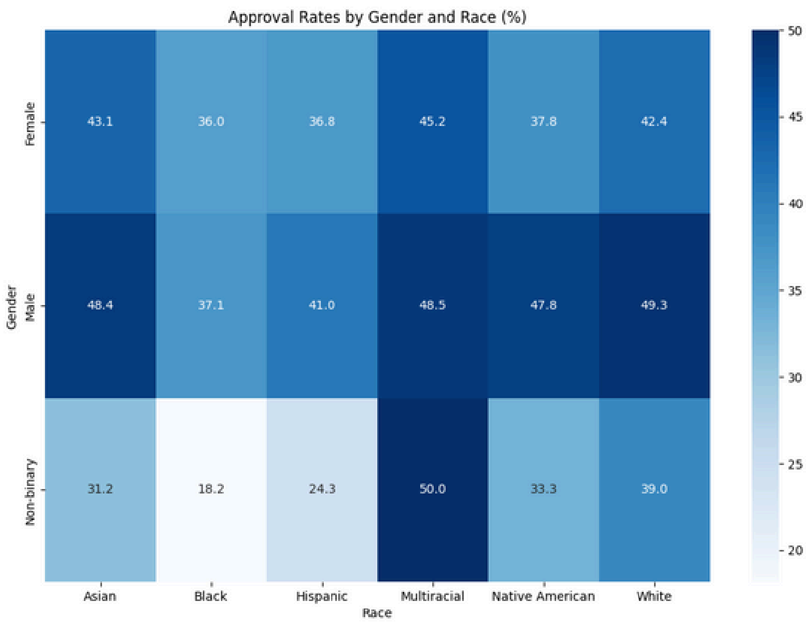
Approval Rates by Gender and Race (%)
Gender
Female
43.1
36.0
36.8
45.2
37.8
42.4
Male
48.4
37.1
41.0
48.5
47.8
49.3
Non-binary
31.2
18.2
24.3
50.0
33.3
39.0
Asian
Black
Hispanic
Multiracial
Native American
White
Race
50
45
40
35
30
25
20

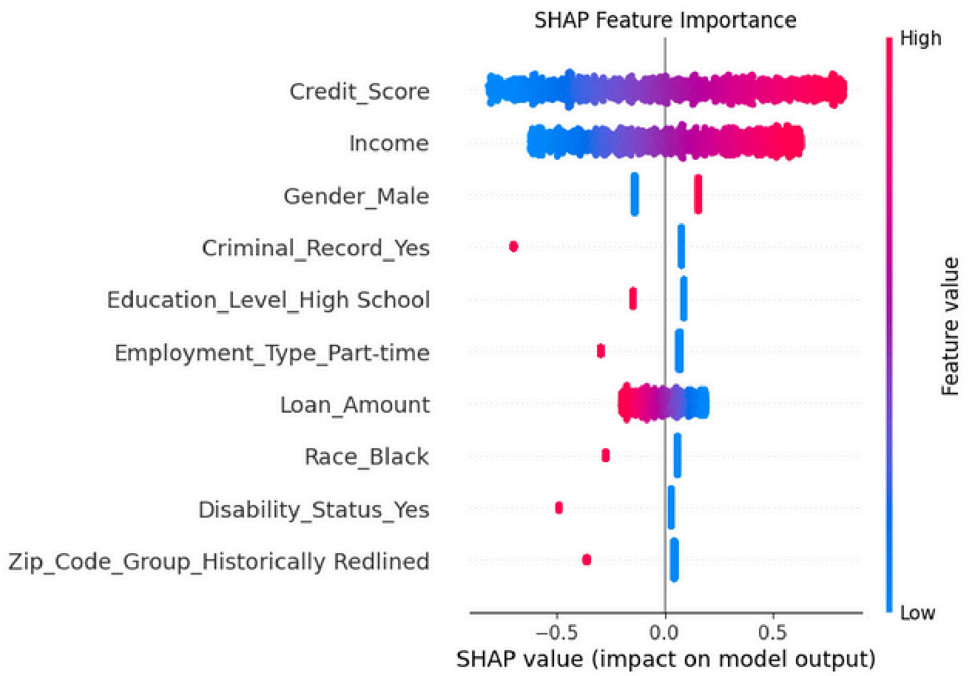
SHAP Feature Importance
Credit_Score
Income
Gender_Male
Criminal_Record_Yes
Education_Level_High School
Employment_Type_Part-time
Loan_Amount
Race_Black
Disability_Status_Yes
Zip_Code_Group_Historically Redlined
−0.5
0.0
0.5
SHAP value (impact on model output)
High
Feature value
Low